

# NLP 2

Inteligencia artificial avanzada para la ciencia  
de datos II Modulo 5 NLP 2

transformers

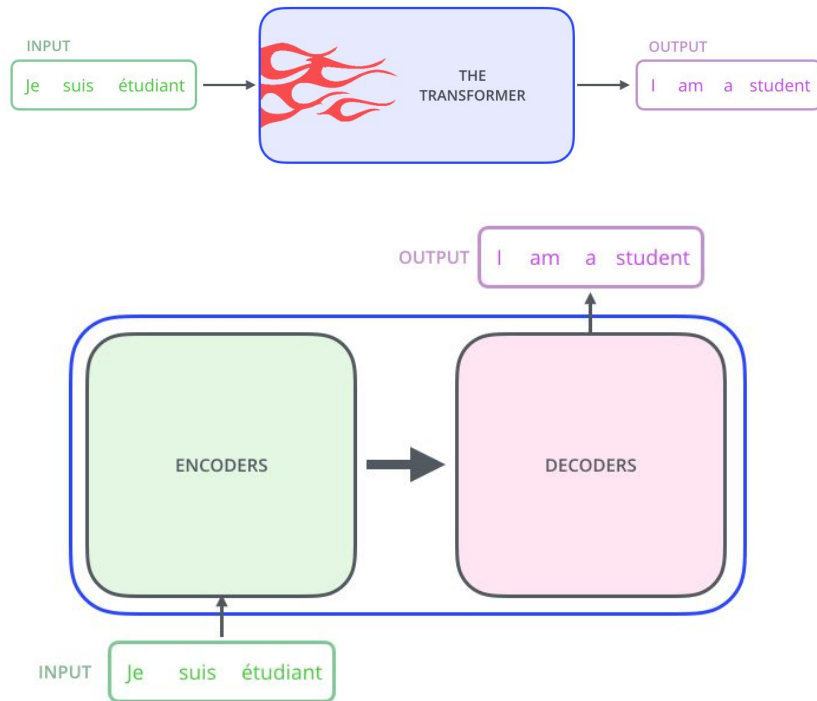
# Transformer

Transformers reuse the model encoder/decoder from seq2seq architecture.

Seq2seq propose two networks (encoding and decoding) and an intermediate state with fix length (vector context) or dynamic (attention layer).

Transformer brings a new paradigm in 2017 with the paper “Attention is all you need” by Google

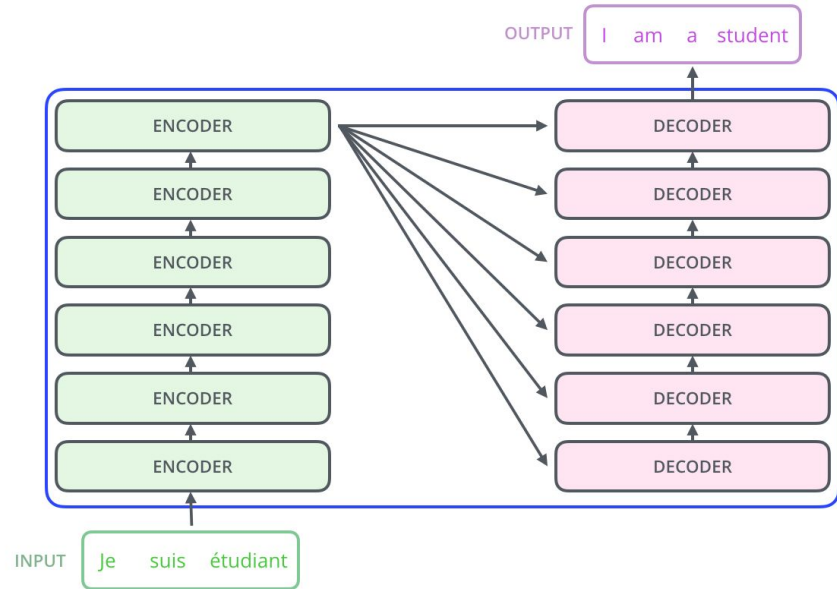
<https://arxiv.org/abs/1706.03762>



# In high level

The high level architecture actually considers **not only one encoder/decoder**, actually there are several encoders/decoders connected by **putting them in a stack**.

The last encoder output will serve as input to the first decoder in the stack.

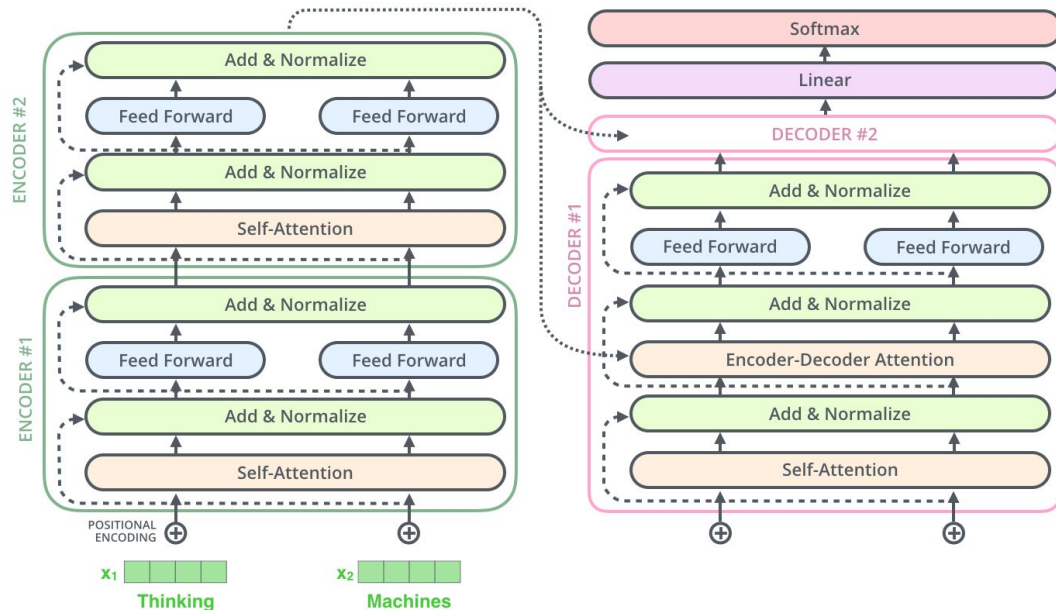


# Low level architecture

Transformer model is quite complex but we will analyze every step.

Some important highlight to remember is the **self-attention** component, but we will go there later in this course.

Lets begin in the first step, you will notice a **positional encoding transformation** just in the very beginning.

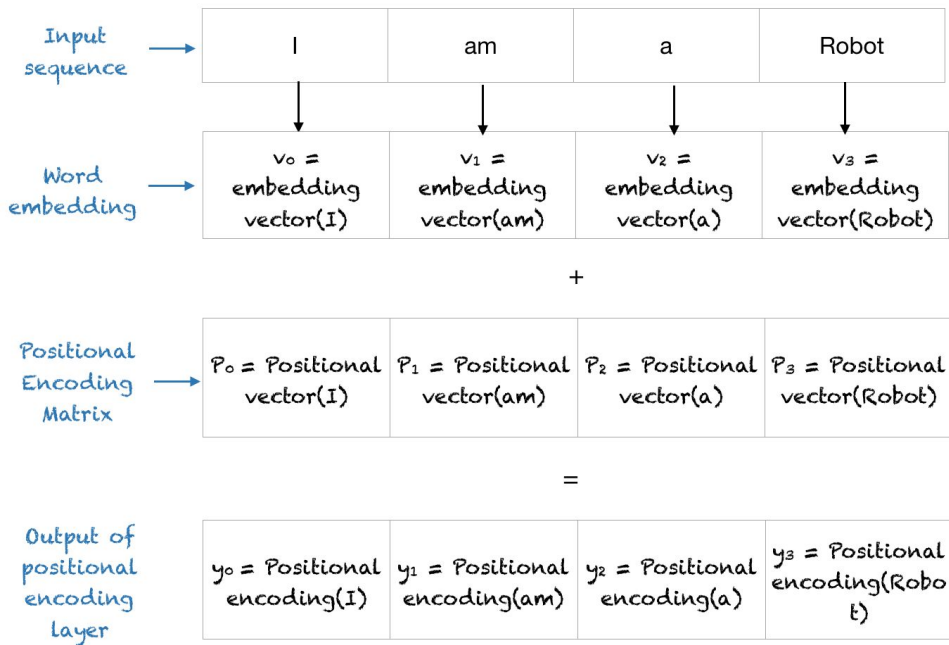


# Positional encoding

A positional encoding is added to the **input embeddings of tokens**.

The purpose of positional encoding is to provide the model with information about the **relative positions of tokens** in the input sequence.

It allows the model to distinguish between tokens based on their position in **addition to their semantic content**.



# How to encode position?

$$P(k, 2i) = \sin\left(\frac{k}{n^{2i/d}}\right)$$

$$P(k, 2i + 1) = \cos\left(\frac{k}{n^{2i/d}}\right)$$

Here:

k: Position of an object in the input sequence,  $0 \leq k < L/2$

d: Dimension of the output embedding space

$P(k, j)$ : Position function for mapping a position k in the input sequence to index (k, j) of the positional matrix

n: User-defined scalar, set to 10,000 by the authors of [Attention Is All You Need](#).

i: Used for mapping to column indices  $0 \leq i < d/2$ , with a single value of i maps to both sine and cosine functions

Sequence	Index of token	Positional Encoding Matrix			
I	0	$P_{00}$	$P_{01}$	...	$P_{0d}$
am	1	$P_{10}$	$P_{11}$	...	$P_{1d}$
a	2	$P_{20}$	$P_{21}$	...	$P_{2d}$
Robot	3	$P_{30}$	$P_{31}$	...	$P_{3d}$

Positional Encoding Matrix for the sequence 'I am a robot'

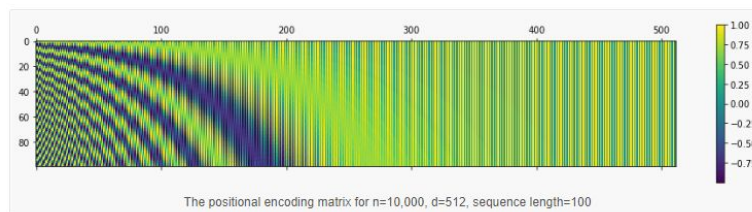
# Example

**Positional encodings** are typically generated using trigonometric functions or other mathematical functions that produce a unique encoding for **each position in the sequence**. These encodings are then added element-wise to the token embeddings.

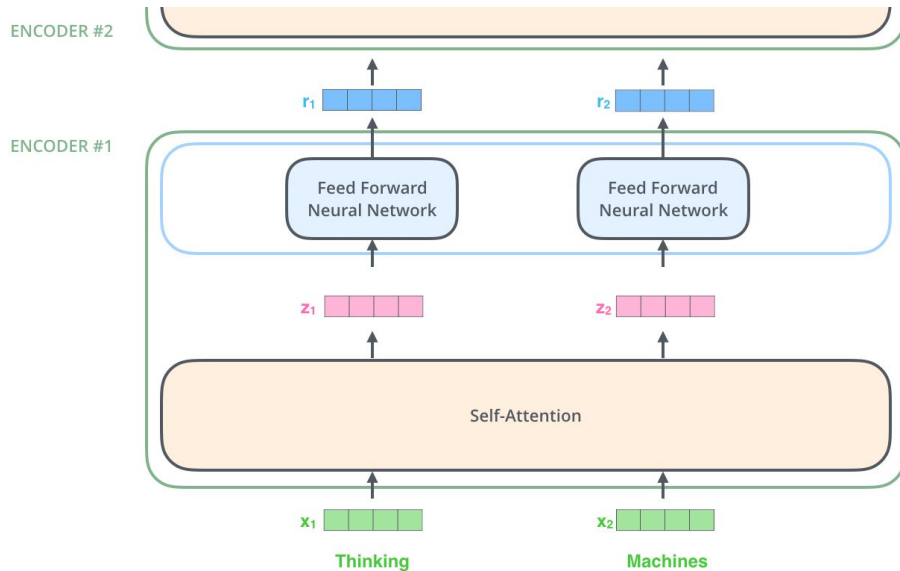
The choice of functions and encoding methods may vary, but one common method is to use a combination of sine and cosine functions with different frequencies to create a distinct encoding for each position

Sequence	Index of token, $k$	Positional Encoding Matrix with $d=4$ , $n=100$			
		$i=0$	$i=0$	$i=1$	$i=1$
I	0	$P_{00}=\sin(0)$ = 0	$P_{01}=\cos(0)$ = 1	$P_{02}=\sin(0)$ = 0	$P_{03}=\cos(0)$ = 1
am	1	$P_{10}=\sin(1/1)$ = 0.84	$P_{11}=\cos(1/1)$ = 0.54	$P_{12}=\sin(1/10)$ = 0.10	$P_{13}=\cos(1/10)$ = 1.0
a	2	$P_{20}=\sin(2/1)$ = 0.91	$P_{21}=\cos(2/1)$ = -0.42	$P_{22}=\sin(2/10)$ = 0.20	$P_{23}=\cos(2/10)$ = 0.98
Robot	3	$P_{30}=\sin(3/1)$ = 0.14	$P_{31}=\cos(3/1)$ = -0.99	$P_{32}=\sin(3/10)$ = 0.30	$P_{33}=\cos(3/10)$ = 0.96

Positional Encoding Matrix for the sequence 'I am a robot'



# Inside encoder



Now that we have our positional embedding, let's move inside each encoder layer.

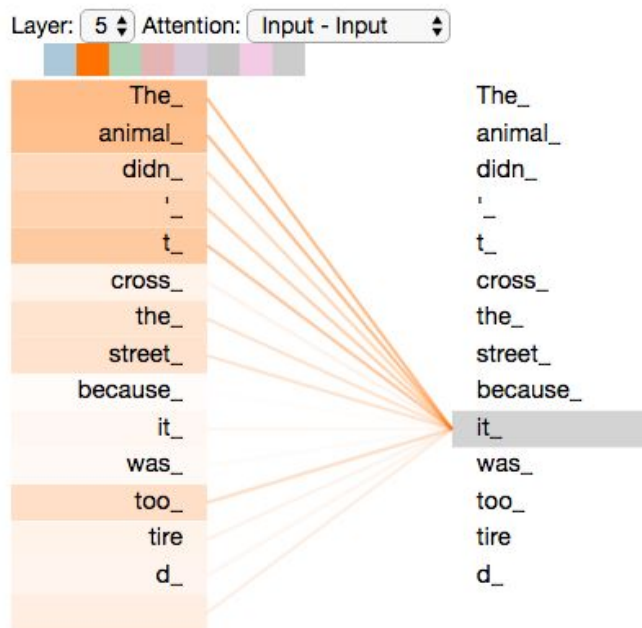
Original paper considers an encoder size of 6, but this is an hyperparameter to explore. (GPT-2, which has 124M parameters, has 12 transformer layers in its architecture.).

Each encoder has two steps

- **Self-attention**
- **FFN**



# Self-attention



As we discuss in last lecture, attention helped seq2seq decoders to produce a new prediction based on the attention step in the input.

Transformers bring self attention as a problem that matters in encoder too.

We need to focus on the relation of the tokens in the input.

For example

"The animal didn't cross the street because it was too tired"

Who was too tired, animal or street?

Self attention brings a better encoding that includes this information

# Queries, keys, values

Let's put it into an intuitive way of think, we need a **final vector that best encode attention**, and we have a **N input set of tokens**.

For every token embedding we need the **"score relation"** with **all the other tokens in the input**.

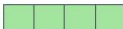
So, for every X vector we will produce three more vectors (usually less in dimension) q,k,v by using three weight matrix ( $W_q, W_k, W_v$ ) in our model.

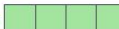
Input

Thinking


Machines


Embedding

$X_1$  


$X_2$  


Queries

$q_1$  


$q_2$  


Keys

$k_1$  

$k_2$  

Values

$v_1$  

$v_2$  



$W^Q$



$W^K$



$W^V$

# Self attention mechanism

This example shows how this 3 vectors helps input “thinking” to be related/scored with all the other vectors in 4 steps.

1. Q1 dot product with all the kn vectors
2. Then it's divided by 8 (the square root of the dimension of the key vectors used in the paper – 64. This leads to having more stable gradients)
3. Softmax normalizes the scores so they're all positive and add up to 1.
4. Multiply value with softmax (intuition here is get a X value transformed since not relevant will be multiplied by near 0 values)
- 5.

Input

Embedding

Queries

Keys

Values

Score

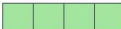
Divide by 8 ( $\sqrt{d_k}$ )

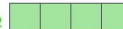
Softmax

Softmax  
X  
Value

Thinking

Machines

$x_1$  

$x_2$  

$q_1$  

$q_2$  

$k_1$  

$k_2$  

$v_1$  

$v_2$  

$q_1 \cdot k_1 = 112$

$q_1 \cdot k_2 = 96$

14

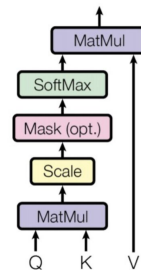
12

0.88

0.12

$v_1$  

$v_2$  



# In matrix operations

The representation of this process in matrix

might look as this:

- X represents all the input vectors ,
- W the matrices for get q,k,x
- Z is the resulting new set of input vectors in a better representation (with attention).

$$\text{softmax} \left( \frac{\begin{matrix} \text{Q} & \text{K}^T \\ \begin{matrix} \text{2x2 grid} \end{matrix} & \begin{matrix} \text{2x2 grid} \end{matrix} \end{matrix} \times \right) \begin{matrix} \text{V} \\ \begin{matrix} \text{2x2 grid} \end{matrix} \end{matrix} \right) = \begin{matrix} \text{Z} \\ \begin{matrix} \text{2x2 grid} \end{matrix} \end{matrix}$$

$$\begin{matrix} \text{X} \\ \begin{matrix} \text{2x4 grid} \end{matrix} \end{matrix} \times \begin{matrix} \text{W}^{\text{Q}} \\ \begin{matrix} \text{4x2 grid} \end{matrix} \end{matrix} = \begin{matrix} \text{Q} \\ \begin{matrix} \text{2x2 grid} \end{matrix} \end{matrix}$$

$$\begin{matrix} \text{X} \\ \begin{matrix} \text{2x4 grid} \end{matrix} \end{matrix} \times \begin{matrix} \text{W}^{\text{K}} \\ \begin{matrix} \text{4x2 grid} \end{matrix} \end{matrix} = \begin{matrix} \text{K} \\ \begin{matrix} \text{2x2 grid} \end{matrix} \end{matrix}$$

$$\begin{matrix} \text{X} \\ \begin{matrix} \text{2x4 grid} \end{matrix} \end{matrix} \times \begin{matrix} \text{W}^{\text{V}} \\ \begin{matrix} \text{4x2 grid} \end{matrix} \end{matrix} = \begin{matrix} \text{V} \\ \begin{matrix} \text{2x2 grid} \end{matrix} \end{matrix}$$

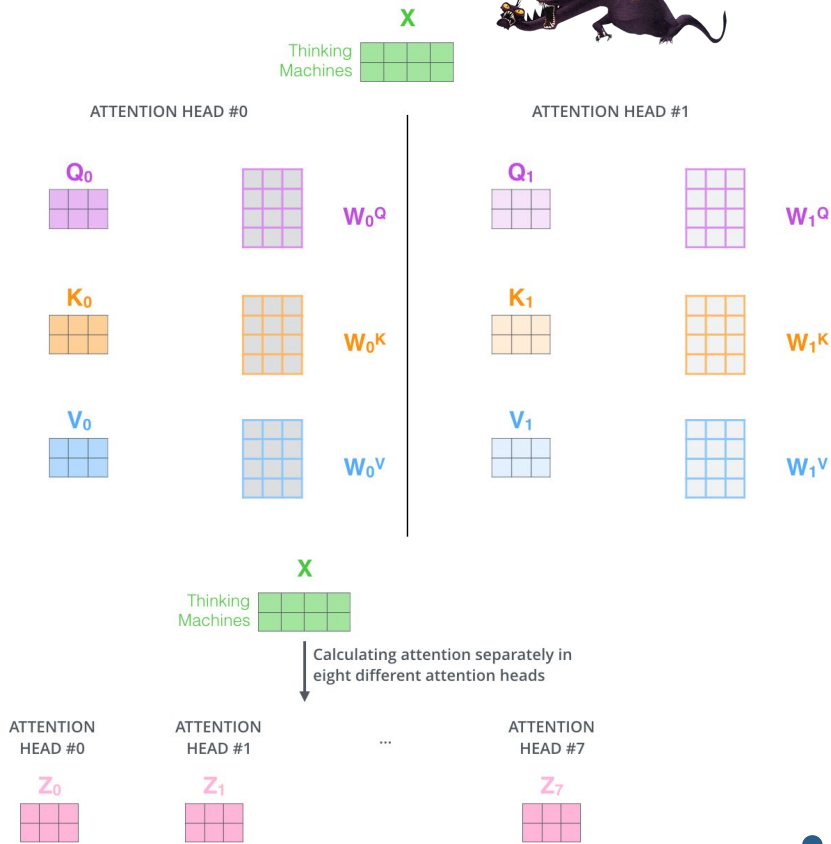
# “multi-headed” attention



Transformers paper brings an optimization, since embeddings is a very high dimensional space we **should train this in many epochs**.

This was a problem in RNN models but now , we can do it in **parallel with the multi-headed mechanism**, this also helps to explore different types of semantic attention (pragmatics analysis).

Original paper propose 8 “heads” (simultaneous matrix weight calculation), however **this is another hyper-parameter** to explore (GPT 3 has 96).



# “multi-headed” issue

Although this is a great idea, the next layer (FFNN) requires one vector per token ...

The answer is to concatenate the output and reduce it by using another  $W_o$  matrix ... **yes another  $W$  matrix.**

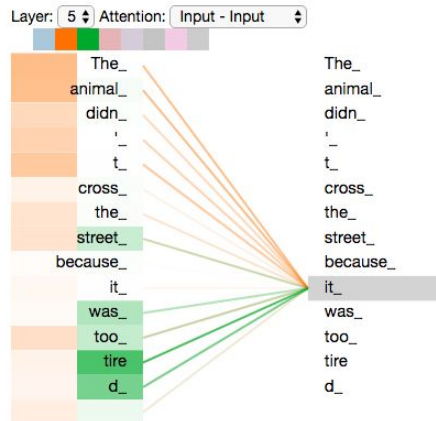
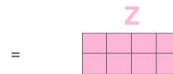
1) Concatenate all the attention heads



2) Multiply with a weight matrix  $W^o$  that was trained jointly with the model

X

3) The result would be the  $Z$  matrix that captures information from all the attention heads. We can send this forward to the FFNN



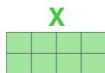
# Self-attention multi-headed recap

We are finishing all the necessary steps to continue to the second part of encoder **FFNN**, lets recap.

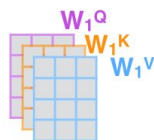
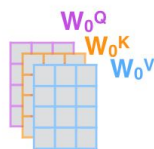
1) This is our input sentence\*

Thinking  
Machines

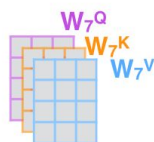
2) We embed each word\*



3) Split into 8 heads.  
We multiply  $X$  or  $R$  with weight matrices



...



4) Calculate attention using the resulting  $Q/K/V$  matrices



...



5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer



...

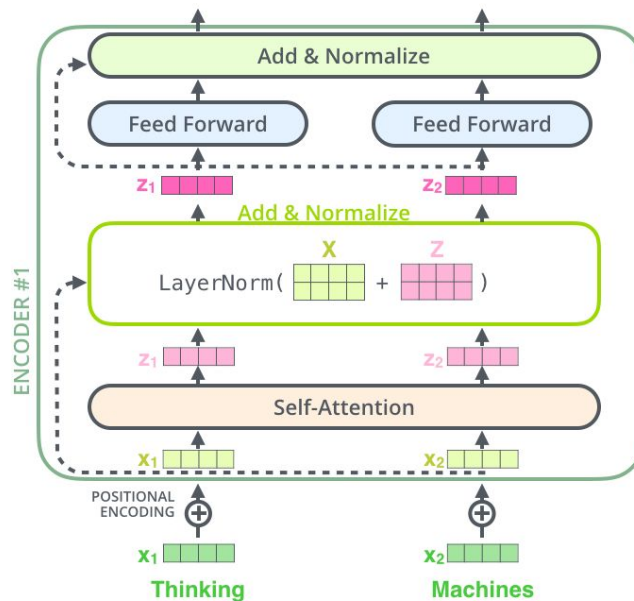


\* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one

# Finishing encoding

In models like the Transformer, which **relies heavily on self-attention** mechanisms, normalization plays a crucial **role in stabilizing the behavior of attention weights**, ensuring that they do not become too large or too small.

Finally FFN plays a crucial role in **extracting, transforming, and generalizing information from the input sequence**. It complements the self-attention mechanism by introducing non-linearity, flexibility, and feature extraction capabilities



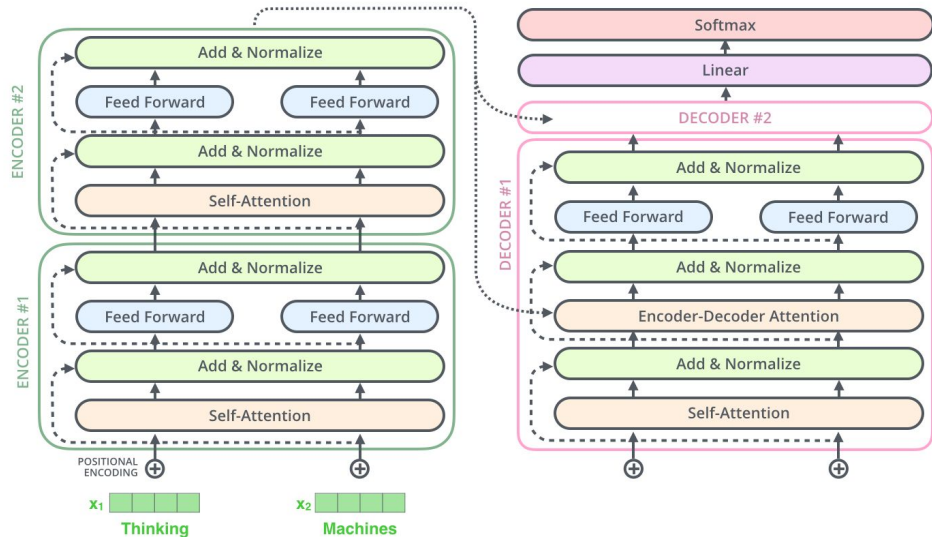
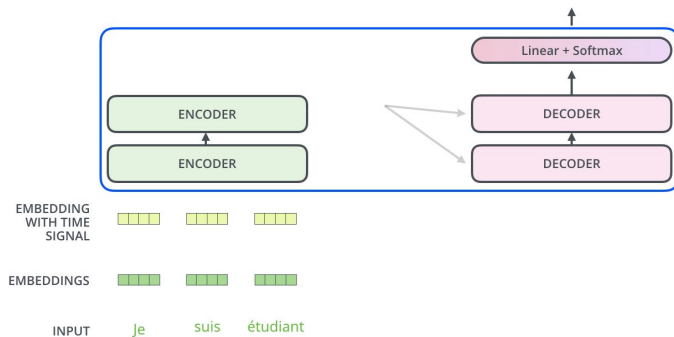


# Starting Decoding

At the end of the top decoder, the output is transformed into a k and v matrices to be use in the “encoder-decoder attention layer”. This process helps the **decoder to focus on the important components of the input (attention)** to predict.

Decoding time step: 1 2 3 4 5 6

OUTPUT





# New vectors to words

Our original embeddings has change to much, in order to convert new vector into tokens we use the last two phases (linear and softmax).

**Linear** is a simple NN to convert vector to a “hot encode” vocabulary representation.

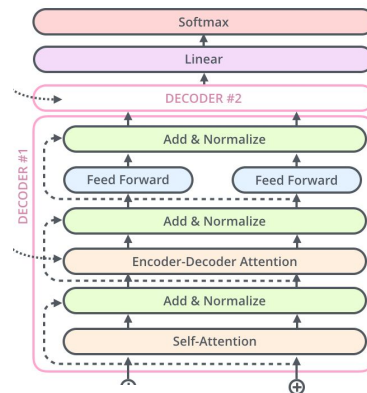
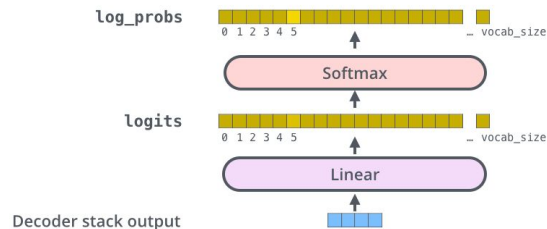
**Softmax** is a function to translate that vector into a probability output.

Which word in our vocabulary is associated with this index?

am

Get the index of the cell with the highest value (argmax)

5



# Training

To train we must include mask examples from our corpus, for example

**“Thanks I am a student”**

Will be trained as

**X1 = “thanks I <mask> a student <EOS>”**

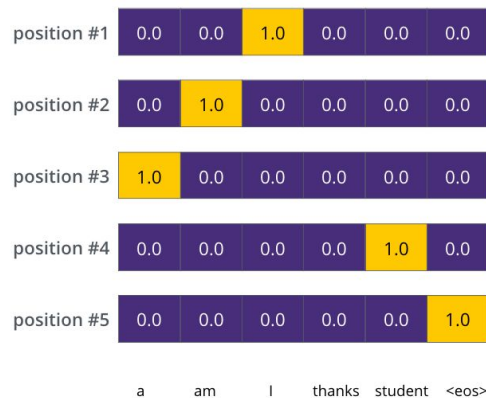
**y1= “am”**

Notice the <mask> input needs to be transformed into am.

The function loss might now measures how well the transformer produces a correct hot encode vector for each position.

## Target Model Outputs

Output Vocabulary: a am I thanks student <eos>



**BERT**, which stands for **Bidirectional Encoder Representations from Transformers**, is a deep learning model that has revolutionized the field of natural language processing (NLP).

Developed by researchers at Google AI in 2018, BERT is a type of transformer model, a neural network architecture that has proven highly effective for a wide range of NLP tasks.

**Try it for free:**

<https://huggingface.co/bert-base-uncased?text=Paris+is+the+%5BMASK%5D+of+France.>

# BERT



## ⚡ Hosted inference API ⓘ

📄 Fill-Mask

Mask token: [MASK]

Paris is the [MASK] of France.

Compute

Computation time on Intel Xeon 3rd Gen Scalable cpu: cached

capital	0.997
heart	0.001
center	0.000
centre	0.000
city	0.000

</> JSON Output

🗲 Maximize

**BERT**, which stands for **Bidirectional Encoder Representations from Transformers**, is a deep learning model that has revolutionized the field of natural language processing (NLP).

Developed by researchers at Google AI in 2018, BERT is a type of transformer model, a neural network architecture that has proven highly effective for a wide range of NLP tasks.

**Try it for free:**

<https://huggingface.co/bert-base-uncased?text=Paris+is+the+%5BMASK%5D+of+France.>

# BERT



## ⚡ Hosted inference API ⓘ

📄 Fill-Mask

Mask token: [MASK]

Paris is the [MASK] of France.

Compute

Computation time on Intel Xeon 3rd Gen Scalable cpu: cached

capital	0.997
heart	0.001
center	0.000
centre	0.000
city	0.000

</> JSON Output

🗲 Maximize

# GPT

**GPT**, which stands for **Generative Pre-trained Transformer**, is a deep learning model that uses transformers just like bert.

Developed by researchers at Open AI in 2018, GTP focus more on decoding task (generative AI) than BERT which focus on understanding better (encoder).

Both has improve in time until reaching both (encoder and decoder) far much better.

**Try it for free:**

<https://huggingface.co/gpt2>

The screenshot shows the Hugging Face GPT-2 demo interface. At the top, it displays 'Downloads last month 25,819,273' with a purple area chart. Below this, there are filters for 'Safetensors', 'Model size 137M params', 'Tensor type F32', and a search icon. The main section is titled 'Hosted inference API' and shows a 'Text Generation' task. A text input box contains the prompt: 'My name is Julien and I like to call myself Julien and you all know me I'm not really good at this but I'd love to put this into perspective so here is not a lot to be talking about and just take my word for'. Below the input box is a 'Compute' button with a 'ctrl+Enter' shortcut. The output area shows the generated text: '0,3'. At the bottom, it indicates 'Computation time on Intel Xeon 3rd Gen Scalable cpu: cached' and a 'JSON Output' toggle. A 'Maximize' button is also visible.

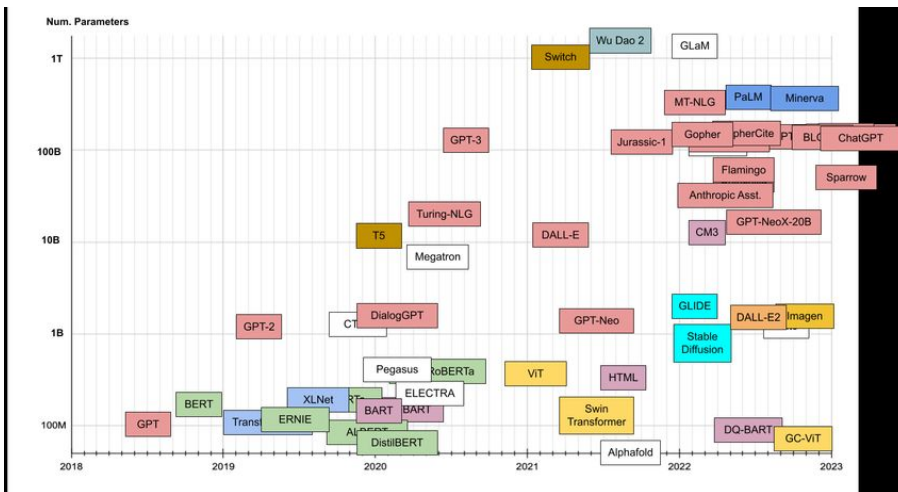
# Transformers timeline

Since BERT and GPT, parameters has been greater and increasing the dataset until today.

Also, more fine tuning models based on top of this models.

This top models are focused on dealing with specific tasks to solve:

- QA
- Creating code
- Solving Maths
- Knowledge retrieve





# Thanks

Do you have any questions?

emmanuel.paez@tec.mx  
Slack #module-5-nlp-1